

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime, timedelta
from IPython.display import display
```

PART 1: BASIC JOINS AND MERGES

These exercises cover fundamental join operations with clear edge cases and detailed explanations.

Key Concepts

1. Merge — `pd.merge()`

Most versatile method for database-style joins

Syntax: `pd.merge(left, right, how='inner', on=None, left_on=None, right_on=None)`

Best for:

- Combining DataFrames based on common columns or indices

Let's create a couple of dataframes to bring together in different ways

```
In [ ]: # @title
# create the first DataFrame
df_mkt = pd.DataFrame({"fruit" : ["apple", "banana", "avocado", "grapes", "blueberry"],
                        "market_price" : [26.99, 29.99, 41.99, 23.99, 26.99] })
print("The first DataFrame")
display(df_mkt)

# create the second DataFrame
df_whl = pd.DataFrame({"fruit" : ["grapes", "blueberries", "banana", "apple", "avocado"],
                        "wholesaler_price" : [19.99, 14.99, 15.99, 20.99, 30.99] })
print("\nThe second DataFrame")
display(df_whl)
```

The first DataFrame

	fruit	market_price
0	apple	26.99
1	banana	29.99
2	avocado	41.99
3	grapes	23.99
4	blueberries	26.99

The second DataFrame

	fruit	wholesaler_price
0	grapes	19.99
1	blueberries	14.99
2	banana	15.99
3	apple	20.99
4	avocado	30.99

```
In [ ]: # joining the DataFrames
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

2. Join — DataFrame.join()

Specialized method for index-based joining

Syntax: `df1.join(df2, how='left', lsuffix='', rsuffix='')`

Best for:

- Combining DataFrames using their indices

```
In [ ]:
```

```
In [ ]:
```

3. Concatenate — pd.concat()

Used for combining DataFrames along an axis

Syntax: `pd.concat([df1, df2], axis=0)`

Best for:

- Stacking DataFrames vertically (axis=0)
- Combining DataFrames side-by-side (axis=1)

```
In [ ]: # Stacking DataFrames vertically (axis=0)
```

```
In [ ]:
```

```
In [ ]: # Combining DataFrames side-by-side (axis=1)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

4. Additional Methods for Combining Data

- `merge_asof()` for time-series-like joining
- `combine_first()` for updating missing values

```
In [ ]: # @title
df_trades = pd.DataFrame({
    'time': pd.to_datetime(['20160525 13:30:00.023',
                           '20160525 13:30:00.038',
                           '20160525 13:30:00.048',
                           '20160525 13:30:00.048',
                           '20160525 13:30:00.048']),
    'ticker': ['MSFT', 'MSFT', 'GOOG', 'GOOG', 'AAPL'],
    'price': [51.95, 51.95, 720.77, 720.92, 98.00],
    'quantity': [75, 155, 100, 100, 100]},
    columns=['time', 'ticker', 'price', 'quantity'])

df_quotes = pd.DataFrame({
    'time': pd.to_datetime(['20160525 13:30:00.023',
                           '20160525 13:30:00.023',
                           '20160525 13:30:00.030',
                           '20160525 13:30:00.041',
                           '20160525 13:30:00.048',
                           '20160525 13:30:00.049',
                           '20160525 13:30:00.072',
                           '20160525 13:30:00.075']),
    'ticker': ['GOOG', 'MSFT', 'MSFT', 'MSFT', 'GOOG', 'AAPL', 'GOOG', 'MSFT'],
    'bid': [720.50, 51.95, 51.97, 51.99, 720.50, 97.99, 720.50, 52.01],
    'ask': [720.93, 51.96, 51.98, 52.00, 720.93, 98.01, 720.88, 52.03]},
    columns=['time', 'ticker', 'bid', 'ask'])
display(df_trades)
display(df_quotes)
```

In []:

In []:

```
In [ ]: # @title
# First DataFrame with some missing values
sales_data = pd.DataFrame({
    'Month': ['January', 'February', 'March', 'April'],
    'Sales': [200, None, 150, None]
})

# Second DataFrame with historical sales data
backup_data = pd.DataFrame({
    'Month': ['January', 'February', 'March', 'April'],
    'Sales': [180, 220, None, 170]
})

display(sales_data)
display(backup_data)
```

```
In [ ]: # Use combine_first to update missing values in sales_data with values from backup_
# Display the updated sales data
```

Task: Basic Join Types

Understanding Different Join Types The hidden code below creates two DataFrames: 'df_students' and 'df_grades'

Key Learning Points:

- INNER join: Only keeps matches present in both DataFrames
- LEFT join: Keeps all records from left DataFrame, fills missing with NaN
- RIGHT join: Keeps all records from right DataFrame, fills missing with NaN
- OUTER join: Keeps all records from both DataFrames, fills missing with NaN

Edge Cases:

- Students with no grades
- Grades with no matching student
- Duplicate student entries

Instructions Create an inner, left, right and outer joins of df_students and df_grades on student_id.

```
In [ ]: # @title
# Sample data
df_students = pd.DataFrame({
    'student_id': [1, 2, 3, 4, 5, 5], # Note: ID 5 is duplicated
```

```

    'name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Eve'],
    'major': ['CS', 'Physics', 'Math', 'CS', 'Physics', 'Physics']
})

df_grades = pd.DataFrame({
    'student_id': [1, 2, 2, 6, 7], # IDs 6,7 don't exist in students
    'course': ['Math', 'Physics', 'CS', 'Math', 'Physics'],
    'grade': [85, 92, 88, 90, 87]
})

```

In []:

Create dataframes based on the different types of joins

Inner, left, right and outer

In []: *# Solutions demonstrating different join types*

```

In [ ]: print("Inner Join")
        display(inner_join)

        print("Outer Join")
        display(outer_join)

        print("Left Join")
        display(left_join)

        print("Right Join")
        display(right_join)

```

Time-Based Joins with Window Functions

```

In [ ]: # Time-Based Joins with Window Functions

        # Sample data
        stock_prices = pd.DataFrame({
            'timestamp': pd.date_range('2024-01-01', periods=5, freq='H'),
            'price': [100, 101, 99, 102, 103]
        })

        company_events = pd.DataFrame({
            'timestamp': pd.date_range('2024-01-01 01:30:00', periods=3, freq='2H'),
            'event_type': ['News', 'Earnings', 'Press Release']
        })

```

PART 2: FASTSHIP, INC CASE

— FastShip Transactional Data —

Shipments:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/shipments.csv>

Orders:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/orders.csv>

Feedback:

https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/cmr_feedback.csv

— SpeedyDelivery Data —

Speedy Orders:

https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/speedy_orders.csv

GPS Tracking:

https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/gps_tracking.csv

Warehouse:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/warehouse.csv>

— Customer Support —

Segments:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/segments.csv>

Interactions:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/interactions.csv>

Outcomes:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/outcomes.csv>

— Product Analytics —

Inventory:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/inventory.csv>

Sales: <https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/sales.csv>

Products:

<https://raw.githubusercontent.com/iamctodd/datasets/refs/heads/main/joins/products.csv>

Part 1: Combining the data

Create a summary report showing:

- Total orders per region
- Average delivery time
- Average customer rating

- Percentage of orders with feedback

```
In [ ]: # First import all the data!
df_orders = pd.read_csv('https://raw.githubusercontent.com/iamctodd/datasets/refs/h
df_ship = pd.read_csv('https://raw.githubusercontent.com/iamctodd/datasets/refs/hea
df_fbk = pd.read_csv('https://raw.githubusercontent.com/iamctodd/datasets/refs/head
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]: # Create a summary report
```

```
In [ ]:
```

```
In [ ]: # Percentage of orders with feedback
```

```
In [ ]: # Total orders per region
""" YOUR CODE HERE """
```

```
In [ ]: # Average delivery times
""" YOUR CODE HERE """
```

```
In [ ]: # Average Customer Rating
""" YOUR CODE HERE """
```

```
In [ ]: # Percent of orders with feedback
""" YOUR CODE HERE """
```

What are the key strategic challenges facing FastShip Inc.?

- ADD YOUR ANSWER HERE

How might an integrated data analytics help address these challenges?

- ADD YOUR ANSWER HERE

Evaluate the timing of the SpeedyDelivery acquisition. What additional complexities does this add to Pablo's analytics integration project

- This is my reply to this quesiont

Write the Python code needed to:

- Calculate delivery times
- Create a unified customer view across both companies

- Calculate standardized delivery times
- Integrate GPS tracking data with delivery performance metrics

```
In [ ]: # Calculate delivery times
        """ YOUR CODE HERE """
```

```
In [ ]: # Create a unified customer view across both companies
        """ YOUR CODE HERE """
```

```
In [ ]: # Calculate standardized delivery times
        """ YOUR CODE HERE """
```

```
In [ ]: # Integrate GPS tracking with delivery performance
        """ YOUR CODE HERE """
```

```
In [ ]:
```

****Evaluate the timing of the SpeedyDelivery acquisition. ****

- Your Answer Here

What additional complexities does this add to Pablo's analytics integration project?

- Your Answer Here
- **Create a customer service report by**
 - Combine all three datasets to analyze resolution times by segment
 - Identify customers with interactions but missing outcomes
 - Calculate the average resolution time by channel and segment

```
In [ ]: """ YOUR CODE HERE """
```

- Create a complete sales report with **** product details and inventory levels **** Analyze sales patterns by category **** Calculate revenue by category and date**

```
In [ ]:
```

- **Address the following:** **** Different column names and metrics across both companies** **** Inconsistent business rules (e.g., delivery time calculations)** **** Real-time GPS data integration** **** Historical data compatibility**