

# Fundamentals of Business Analytics - Week 10 Session 1

## Notes

- Linear regression is a statistical method used to model the relationship between a dependent variable and one or more independent variables by fitting a linear equation to the data.
- A scatter plot can be used to:
  - Visualize the relationship between X and Y variables.
- Only one independent variable, X.
- Relationship between X and Y is described by a linear function.
- Changes in Y are assumed to be related to changes in X.

### Simple Linear Regression Model

- $Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$
- Dependent Variable = Population Y intercept + (Population Slope Coefficient \* Independent Variable) + Random Error term
- The simple linear regression equation provides an estimate of the population regression line.
- $b_0$  is the estimated mean value of Y when the value of X is zero.
- $b_1$  is the estimated change in the mean value of Y as a result of a one-unit increase in X.
  - homoscedasciticty

In [ ]:

## House Price Prediction Using Simple Linear Regression

A real estate agent wishes to examine the relationship between the selling price of a home and its size (measured in square feet). This notebook reads house price data from a CSV file,

visualizes it with a scatter plot, and performs simple linear regression to model the relationship between house size and price.

```
In [1]: # IMPORT OUR LIBRARIES
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.formula.api as smf
from statsmodels.stats.anova import anova_lm

# plotting defaults
sns.set(style='whitegrid')
%matplotlib inline
```

## STEP 1 - SET OBJECTIVE

**t test for a population slope:**

Is there a linear relationship between House prices and Size?

Null and alternative hypotheses:

H<sub>0</sub>:  $\beta_1 = 0$  (no linear relationship)

H<sub>1</sub>:  $\beta_1 \neq 0$  (linear relationship does exist)

Significance level: 0.05 (95% CI)

```
In [2]: # Main analysis: Load data, visualize, fit OLS, report t-test and 95% CI for slope
df = pd.read_csv('https://raw.githubusercontent.com/Kartavya-Jharwal/Kartavya_Busin
print("Loaded 'house_prices.csv'.")
```

```
# quick head
print(df.head())
```

Loaded 'house\_prices.csv'.

	price	area
0	245	1400
1	312	1600
2	279	1700
3	308	1875
4	199	1100

## Scatter Plot

```
In [3]: # Create a scatter plot
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x="area", y="price", s=100, color='navy')
plt.title("House price model: Scatter Plot", fontsize=16)
plt.xlabel("Square Feet", fontsize=12)
```

```
plt.ylabel("House Price ($1000s)", fontsize=12)
plt.grid(False)
plt.xlim(0)
plt.ylim(0)
plt.tight_layout()
plt.show()
```



## RUN THE LINEAR REGRESSION

```
In [4]: # Linear regression
X = df["area"]
y = df["price"]
# Use formula-based interface instead of sm.OLS
model = smf.ols('price ~ area', data=df).fit()
print('\n==== OLS Summary ===')
print(model.summary())
```

```

    === OLS Summary ===
                    OLS Regression Results
    =====
Dep. Variable:           price   R-squared:      0.581
Model:                 OLS     Adj. R-squared:  0.528
Method:                Least Squares F-statistic:   11.08
Date: Mon, 10 Nov 2025   Prob (F-statistic): 0.0104
Time: 14:50:12          Log-Likelihood: -50.290
No. Observations:      10     AIC:             104.6
Df Residuals:          8     BIC:             105.2
Df Model:              1
Covariance Type:       nonrobust
=====
            coef    std err      t      P>|t|      [0.025      0.975]
-----
Intercept    98.2483    58.033    1.693    0.129    -35.577    232.074
area         0.1098    0.033    3.329    0.010      0.034     0.186
=====
Omnibus:                  1.066 Durbin-Watson:        3.222
Prob(Omnibus):            0.587 Jarque-Bera (JB):  0.779
Skew:                      0.399 Prob(JB):        0.677
Kurtosis:                  1.890 Cond. No.    7.82e+03
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 7.82e+03. This might indicate that there are strong multicollinearity or other numerical problems.

## Check Residual Assumptions

### NORMALITY OF ERROR

```
In [5]: from scipy.stats import shapiro

# Get the residuals from the model
residuals = model.resid

# Perform the Shapiro-Wilk test for normality
shapiro_test = shapiro(residuals)

# Print the results
print(f"Shapiro-Wilk Test Statistic: {shapiro_test.statistic:.4f}")
print(f"Shapiro-Wilk p-value: {shapiro_test.pvalue:.4f}")

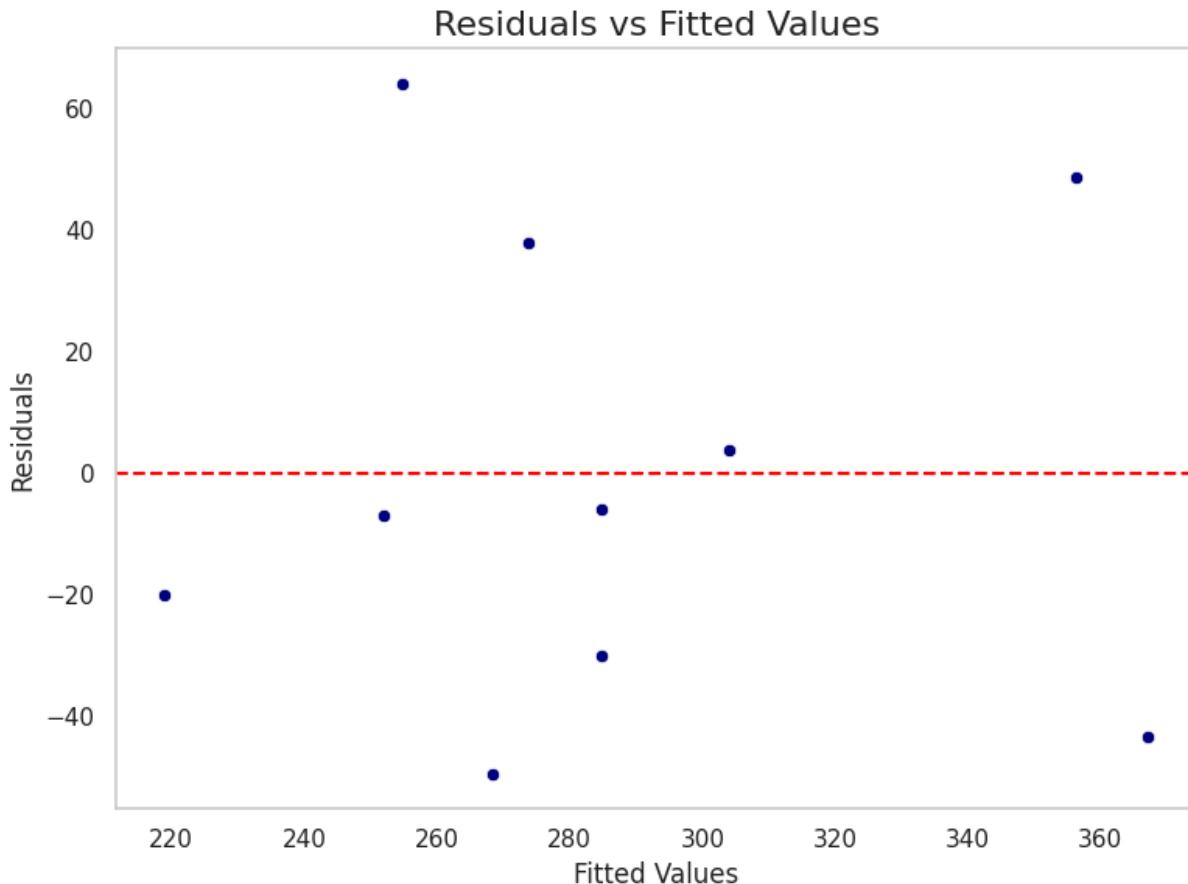
# Interpret the results
alpha = 0.05
if shapiro_test.pvalue > alpha:
    print("The residuals appear to be normally distributed (fail to reject H0)")
else:
    print("The residuals do not appear to be normally distributed (reject H0)")
```

```
Shapiro-Wilk Test Statistic: 0.9353
Shapiro-Wilk p-value: 0.5017
The residuals appear to be normally distributed (fail to reject H0)
```

#### EQUAL VARIANCE CHECK

```
In [6]: import statsmodels.stats.api as sms

# Get fitted values
fitted_values = model.fittedvalues
# Create a scatter plot of residuals vs fitted values
plt.figure(figsize=(8, 6))
sns.scatterplot(x=fitted_values, y=residuals, color='navy')
plt.axhline(y=0, color='red', linestyle='--')
plt.title("Residuals vs Fitted Values", fontsize=16)
plt.xlabel("Fitted Values", fontsize=12)
plt.ylabel("Residuals", fontsize=12)
plt.grid(False)
plt.tight_layout()
plt.show()
```



## Simple Linear Regression Prediction

Making a prediction of house price based on area in square feet. We have fitted a simple linear regression model to predict house price based on square footage using the equation:

$y^{\wedge}$  = Predicted house price (in \$1000s)

---

### Prediction Example

To predict the price of a house with 2000 square feet:

$$y^{\wedge} = 98.24833 + 0.10977 \cdot 2000$$

Predicted House Price: \$317,790  $y^{\wedge} = 98.24833 + 219.54 = 317.79$

The python linear model object contains a .predict() method, predictions can be generated by calling this method and providing the input values in a DataFrame the names of the variables must match the model

---

```
model.predict(pd.DataFrame({'area': [2000]}))
```

```
In [7]: # Get the coefficients from the trained model
intercept = model.params['Intercept']
slope = model.params['area']

# Prompt the user for an area input
try:
    input_area = float(input("Enter the house area in square feet (e.g., 1500): "))

    # Use the model's predict method with a DataFrame input
    predicted_price = model.predict(pd.DataFrame({'area': [input_area]}))[0]

    # Pretty print the results
    print(f"\n--- Prediction Result ---")
    print(f"Input Area: {input_area:.0f} sq ft")
    print(f"Predicted Price: ${predicted_price:.2f} (in $1000s)")
    print(f"Using formula: price = {intercept:.4f} + {slope:.4f} * area")
except ValueError:
    print("Invalid input. Please enter a numerical value for the area.")
```


Enter the house area in square feet (e.g., 1500): 2000

```
--- Prediction Result ---
Input Area: 2,000 sq ft
Predicted Price: $317.78 (in $1000s)
Using formula: price = 98.2483 + 0.1098 * area
```

## Plot R2

```
In [8]: # plot regression line
plt.figure(figsize=(8,6))
sns.regplot(x='area', y='price', data=df, ci=95, scatter_kws={'s':40, 'alpha':0.6})
plt.title('House Price vs area with OLS fit')
plt.xlabel('area (sq ft)')
```

```
plt.ylabel('Price (currency units)')  
plt.show()
```



## Hypothesis Test for the Slope Using the Student's t-Distribution

To test whether the independent variable (e.g., area) significantly predicts the dependent variable (e.g., price), we perform a t-test on the slope coefficient using the Student's t-distribution.

Hypotheses: Null Hypothesis (no relationship between area and price)  $H_0: \beta_1 = 0$

Alternative Hypothesis (a significant linear relationship exists)

$H_1: \beta_1 \neq 0$

The t-statistic measures how many standard errors the estimated slope is away from zero

```
In [9]: # t-test for slope (area)  
slope_t = model.tvalues['area']  
slope_p = model.pvalues['area']  
ci = model.conf_int(alpha=0.05).loc['area']
```

```
print(f"\nSlope t-statistic: {slope_t:.4f}")
print(f"Slope p-value: {slope_p:.4e}")
print(f"95% CI for slope: [{ci[0]:.4f}, {ci[1]:.4f}]")

alpha = 0.05
if slope_p < alpha:
    print('\nConclusion: Reject H0 – there is evidence of a linear relationship bet
else:
    print('\nConclusion: Fail to reject H0 – no evidence of linear relationship at
```

Slope t-statistic: 3.3294  
Slope p-value: 1.0394e-02  
95% CI for slope: [0.0337, 0.1858]

Conclusion: Reject H0 – there is evidence of a linear relationship between area and Price at alpha=0.05.