

## Task B: Car Emissions Analysis

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from scipy.stats import shapiro
```

```
In [2]: from google.colab import files
import io

print("Please upload your CSV or Excel file.")
uploaded = files.upload()

# Assuming only one file is uploaded
for fn in uploaded.keys():
    print(f'User uploaded file "{fn}" with length {len(uploaded[fn])} bytes')

    if fn.endswith('.csv'):
        df = pd.read_csv(io.StringIO(uploaded[fn].decode('utf-8')))
    elif fn.endswith('.xls', '.xlsx'):
        df = pd.read_excel(io.BytesIO(uploaded[fn]))
    else:
        raise ValueError("Unsupported file format. Please upload a .csv or .xlsx file")

# Display basic information
print("Dataset loaded successfully!")
print(f"Shape: {df.shape}")
print("\nFirst few rows:")
print(df.head())
print("\nColumn names:")
print(df.columns.tolist())
print("\nSummary statistics:")
print(df.describe())
```

Please upload your CSV or Excel file.

Choose Files No file chosen

Upload widget is only available when the cell has

been executed in the current browser session. Please rerun this cell to enable.

```
Saving Cars.xlsx to Cars.xlsx
User uploaded file "Cars.xlsx" with length 5961 bytes
Dataset loaded successfully!
Shape: (36, 5)
```

First few rows:

	Car	Model	Volume	Weight	CO2
0	Toyota	Aygo	1000	790	99
1	Mitsubishi	Space Star	1200	1160	95
2	Skoda	Citigo	1000	929	95
3	Fiat	500	900	865	90
4	Mini	Cooper	1500	1140	105

Column names:

```
['Car', 'Model', 'Volume', 'Weight', 'CO2']
```

Summary statistics:

	Volume	Weight	CO2
count	36.000000	36.000000	36.000000
mean	1611.111111	1292.277778	102.027778
std	388.975047	242.123889	7.454571
min	900.000000	790.000000	90.000000
25%	1475.000000	1117.250000	97.750000
50%	1600.000000	1329.000000	99.000000
75%	2000.000000	1418.250000	105.000000
max	2500.000000	1746.000000	120.000000

## Dataset Overview

- **Y (Dependent variable):** CO2 emissions (g/km).
- **X<sub>1</sub>:** Engine Volume (cc).
- **X<sub>2</sub>:** Vehicle Weight (kg).

We inspect descriptive statistics and correlations to understand initial relationships and check for multicollinearity (Part 0).

```
In [3]: # Descriptive Statistics
summary_stats = df[['CO2', 'Volume', 'Weight']].describe().T
print('*'*80)
print('DESCRIPTIVE STATISTICS')
print('*'*80)
print(summary_stats)
print('\n' + '*'*80)
print('DATA DISTRIBUTION SUMMARY')
print('*'*80)
print(f"Sample Size: {len(df)}")
print(f"\nCO2 Emissions (g/km): Mean = {df['CO2'].mean():.2f}, Std = {df['CO2'].std:.2f}")
print(f"Engine Volume (cc): Mean = {df['Volume'].mean():.2f}, Std = {df['Volume'].std:.2f}")
print(f"Vehicle Weight (kg): Mean = {df['Weight'].mean():.2f}, Std = {df['Weight'].std:.2f}")
```

```
=====
DESCRIPTIVE STATISTICS
=====

      count      mean       std     min    25%    50%    75%  \
CO2      36.0  102.027778   7.454571  90.0   97.75  99.0  105.00
Volume   36.0  1611.111111  388.975047 900.0  1475.00 1600.0  2000.00
Weight   36.0  1292.277778 242.123889 790.0  1117.25 1329.0  1418.25

      max
CO2      120.0
Volume  2500.0
Weight   1746.0

=====
DATA DISTRIBUTION SUMMARY
=====

Sample Size: 36

CO2 Emissions (g/km): Mean = 102.03, Std = 7.45
Engine Volume (cc): Mean = 1611.11, Std = 388.98
Vehicle Weight (kg): Mean = 1292.28, Std = 242.12
```

## Exploratory Data Analysis (EDA)

### Visual Inspection of Distributions

```
In [4]: # Distribution plots for all variables
fig, axes = plt.subplots(1, 3, figsize=(16, 4))

# CO2 Distribution
axes[0].hist(df['CO2'], bins=25, edgecolor='black', alpha=0.7, color='coral')
axes[0].axvline(df['CO2'].mean(), color='red', linestyle='--', linewidth=2, label=f'CO2 Mean: {df["CO2"].mean():.2f}')
axes[0].axvline(df['CO2'].median(), color='blue', linestyle='--', linewidth=2, label=f'CO2 Median: {df["CO2"].median():.2f}')
axes[0].set_title('CO2 Emissions Distribution', fontweight='bold', fontsize=12)
axes[0].set_xlabel('CO2 (g/km)')
axes[0].set_ylabel('Frequency')
axes[0].legend()
axes[0].grid(alpha=0.3)

# Volume Distribution
axes[1].hist(df['Volume'], bins=25, edgecolor='black', alpha=0.7, color='skyblue')
axes[1].axvline(df['Volume'].mean(), color='red', linestyle='--', linewidth=2, label=f'Volume Mean: {df["Volume"].mean():.2f}')
axes[1].axvline(df['Volume'].median(), color='blue', linestyle='--', linewidth=2, label=f'Volume Median: {df["Volume"].median():.2f}')
axes[1].set_title('Engine Volume Distribution', fontweight='bold', fontsize=12)
axes[1].set_xlabel('Volume (cc)')
axes[1].set_ylabel('Frequency')
axes[1].legend()
axes[1].grid(alpha=0.3)


# Weight Distribution
axes[2].hist(df['Weight'], bins=25, edgecolor='black', alpha=0.7, color='lightgreen')
axes[2].axvline(df['Weight'].mean(), color='red', linestyle='--', linewidth=2, label=f'Weight Mean: {df["Weight"].mean():.2f}')
axes[2].axvline(df['Weight'].median(), color='blue', linestyle='--', linewidth=2, label=f'Weight Median: {df["Weight"].median():.2f}')
axes[2].set_title('Vehicle Weight Distribution', fontweight='bold', fontsize=12)
```

```

axes[2].set_xlabel('Weight (kg)')
axes[2].set_ylabel('Frequency')
axes[2].legend()
axes[2].grid(alpha=0.3)

plt.tight_layout()
plt.show()

```



## Correlation Analysis

```

In [5]: # Correlation matrix
corr_matrix = df[['CO2', 'Volume', 'Weight']].corr()

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Heatmap
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0,
             square=True, linewidths=1, cbar_kws={"shrink": 0.8},
             fmt='.3f', ax=axes[0], vmin=-1, vmax=1)
axes[0].set_title('Correlation Matrix Heatmap', fontweight='bold', fontsize=12)


# Correlation bar plot
corr_with_co2 = corr_matrix['CO2'].drop('CO2')
colors = ['green' if x > 0 else 'red' for x in corr_with_co2.values]
axes[1].barh(corr_with_co2.index, corr_with_co2.values, color=colors, alpha=0.7, ed
axes[1].axvline(0, color='black', linewidth=0.8)
axes[1].set_xlabel('Correlation Coefficient', fontweight='bold')
axes[1].set_title('Correlation with CO2 Emissions', fontweight='bold', fontsize=12)
axes[1].grid(alpha=0.3, axis='x')

for i, v in enumerate(corr_with_co2.values):
    axes[1].text(v + 0.02 if v > 0 else v - 0.02, i, f'{v:.3f}', va='center', ha='left' if v > 0 else 'right', fontweight='bold')

plt.tight_layout()
plt.show()

print('\nCorrelation Interpretation:')
print(f"- Volume-CO2 correlation: {corr_matrix.loc['Volume', 'CO2']:.3f}")
print(f"- Weight-CO2 correlation: {corr_matrix.loc['Weight', 'CO2']:.3f}")
print(f"- Volume-Weight correlation: {corr_matrix.loc['Volume', 'Weight']:.3f} (mul

```



Correlation Interpretation:

- Volume-CO2 correlation: 0.592
- Weight-CO2 correlation: 0.552
- Volume-Weight correlation: 0.754 (multicollinearity check)


## Linearity Check: Scatter Plots with Regression Lines

```
In [6]: # Scatter plots with regression lines - axes starting at (0,0)
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# CO2 vs Volume
axes[0].scatter(df['Volume'], df['CO2'], alpha=0.6, s=50, color='steelblue', edgecolors='black')
z = np.polyfit(df['Volume'], df['CO2'], 1)
p = np.poly1d(z)
axes[0].plot(df['Volume'], p(df['Volume']), "r--", linewidth=2, label=f'y = {z[0]}:{z[1]}x')
axes[0].set_xlim(0, df['Volume'].max() * 1.1)
axes[0].set_ylimit(0, df['CO2'].max() * 1.1)
axes[0].set_xlabel('Engine Volume (cc)', fontweight='bold', fontsize=11)
axes[0].set_ylabel('CO2 Emissions (g/km)', fontweight='bold', fontsize=11)
axes[0].set_title('CO2 vs Engine Volume', fontweight='bold', fontsize=12)
axes[0].legend(loc='upper left')
axes[0].grid(alpha=0.3)

# CO2 vs Weight
axes[1].scatter(df['Weight'], df['CO2'], alpha=0.6, s=50, color='coral', edgecolors='black')
z = np.polyfit(df['Weight'], df['CO2'], 1)
p = np.poly1d(z)
axes[1].plot(df['Weight'], p(df['Weight']), "r--", linewidth=2, label=f'y = {z[0]}:{z[1]}x')
axes[1].set_xlim(0, df['Weight'].max() * 1.1)
axes[1].set_ylimit(0, df['CO2'].max() * 1.1)
axes[1].set_xlabel('Vehicle Weight (kg)', fontweight='bold', fontsize=11)
axes[1].set_ylabel('CO2 Emissions (g/km)', fontweight='bold', fontsize=11)
axes[1].set_title('CO2 vs Vehicle Weight', fontweight='bold', fontsize=12)
axes[1].legend(loc='upper left')
axes[1].grid(alpha=0.3)

plt.tight_layout()
plt.show()
```



## Part 1: Build the Multiple Regression Model

### Regression Model Specification

#### Model Equation:

$$CO_2 = \beta_0 + \beta_1 \times Volume + \beta_2 \times Weight + \epsilon$$

Where:

- **CO<sub>2</sub>** = Carbon dioxide emissions (g/km) - *Dependent Variable (Y)*
- **Volume** = Engine volume (cc) - *Independent Variable (X<sub>1</sub>)*
- **Weight** = Vehicle weight (kg) - *Independent Variable (X<sub>2</sub>)*
- **$\beta_0$**  = Intercept (baseline CO<sub>2</sub> when Volume and Weight are zero)
- **$\beta_1$**  = Coefficient for Volume (effect per cc increase)
- **$\beta_2$**  = Coefficient for Weight (effect per kg increase)
- **$\epsilon$**  = Error term (residuals)

We use **Ordinary Least Squares (OLS)** estimation to obtain coefficients, standard errors, t-statistics, p-values, and confidence intervals.

```
In [7]: X = df[['Volume', 'Weight']]
X = sm.add_constant(X) # adds intercept term
Y = df['CO2']

model = sm.OLS(Y, X).fit()
print('*'*80)
print('MULTIPLE REGRESSION MODEL SUMMARY')
print('*'*80)
print(model.summary())

print('KEY REGRESSION STATISTICS')
print('*'*80)
print(f'R-squared: {model.rsquared:.4f}')
print(f'Adjusted R-squared: {model.rsquared_adj:.4f}')
print(f'F-statistic: {model.fvalue:.4f} (p-value: {model.f_pvalue:.4f})')
print(f'AIC: {model.aic:.2f}, BIC: {model.bic:.2f}'')
```

```
=====
MULTIPLE REGRESSION MODEL SUMMARY
=====

                    OLS Regression Results
=====

Dep. Variable:          CO2    R-squared:         0.377
Model:                 OLS    Adj. R-squared:      0.339
Method:                Least Squares   F-statistic:       9.966
Date:      Wed, 19 Nov 2025   Prob (F-statistic): 0.000411
Time:      14:48:57        Log-Likelihood:     -114.39
No. Observations:      36    AIC:                  234.8
Df Residuals:          33    BIC:                  239.5
Df Model:               2
Covariance Type:       nonrobust
=====

            coef    std err      t      P>|t|      [0.025      0.975]
-----
const    79.6947    5.564    14.322    0.000     68.374    91.016
Volume   0.0078    0.004     1.948    0.060    -0.000     0.016
Weight   0.0076    0.006     1.173    0.249    -0.006     0.021
=====
Omnibus:             4.957    Durbin-Watson:      0.944
Prob(Omnibus):        0.084    Jarque-Bera (JB):  1.836
Skew:                 -0.025    Prob(JB):           0.399
Kurtosis:              1.895    Cond. No.        1.16e+04
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.16e+04. This might indicate that there are strong multicollinearity or other numerical problems.

#### KEY REGRESSION STATISTICS

```
R-squared: 0.3766
Adjusted R-squared: 0.3388
F-statistic: 9.9659 (p-value: 0.0004)
AIC: 234.78, BIC: 239.53
```

## Coefficient checking

```
In [8]: print('*'*80)
print('COEFFICIENTS WITH 95% CONFIDENCE INTERVALS')
print('*'*80)
conf_int = model.conf_int(alpha=0.05)
for var in model.params.index:
    coef = model.params[var]
    se = model.bse[var]
    t_stat = model.tvalues[var]
    p_val = model.pvalues[var]
    ci_low, ci_high = conf_int.loc[var]
    direction = 'positive' if coef > 0 else 'negative'
    decision = 'Reject H0' if p_val < 0.05 else 'Fail to reject H0'
    print(f'Variable: {var}')
```

```
=====
```

COEFFICIENTS WITH 95% CONFIDENCE INTERVALS

```
=====
```

Variable: const  
Variable: Volume  
Variable: Weight

## Hypothesis Testing

### Null and Alternative Hypotheses for Regression Coefficients

#### For Engine Volume ( $\beta_1$ ):

- **$H_0$  (Null Hypothesis):**  $\beta_1 = 0$

*Engine volume has no effect on CO2 emissions (holding weight constant)*

- **$H_1$  (Alternative Hypothesis):**  $\beta_1 \neq 0$

*Engine volume has a significant effect on CO2 emissions (holding weight constant)*

#### For Vehicle Weight ( $\beta_2$ ):

- **$H_0$  (Null Hypothesis):**  $\beta_2 = 0$

*Vehicle weight has no effect on CO2 emissions (holding volume constant)*

- **$H_1$  (Alternative Hypothesis):**  $\beta_2 \neq 0$

*Vehicle weight has a significant effect on CO2 emissions (holding volume constant)*

#### Overall Model Test (F-test):

- **$H_0$  (Null Hypothesis):**  $\beta_1 = \beta_2 = 0$

*Neither predictor has an effect on CO2 emissions (the model is not significant)*

- **$H_1$  (Alternative Hypothesis):** At least one  $\beta \neq 0$

*At least one predictor has a significant effect on CO2 emissions*

**Significance Level:**  $\alpha = 0.05$

**Decision Rule:** Reject  $H_0$  if p-value < 0.05

```
In [9]: print('*80)
print('STATISTICAL HYPOTHESIS TESTS')
print('*80)
print('Overall F-test: H0 -> β_volume = β_weight = 0')
print(f' F-statistic p-value: {model.f_pvalue:.4f} (Decision: {"Reject H0" if mode
print('\nIndividual t-tests: H0 -> β = 0')
for var in ['Volume', 'Weight']:
    coef = model.params[var]
    p_val = model.pvalues[var]
    direction = 'positive' if coef > 0 else 'negative'
    decision = 'Reject H0' if p_val < 0.05 else 'Fail to reject H0'
```

```

    print(f' - {var}: Coef={coef:.4f}, p-value={p_val:.4f}, Direction={direction},

# Variance Inflation Factor (multicollinearity)
vif_data = pd.DataFrame({
    'Feature': X.columns,
    'VIF': [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
})
print('\nVIF Results:')
print(vif_data)

=====
STATISTICAL HYPOTHESIS TESTS
=====

Overall F-test: H0 -> β_volume = β_weight = 0
F-statistic p-value: 0.0004 (Decision: Reject H0)

Individual t-tests: H0 -> β = 0
- Volume: Coef=0.0078, p-value=0.0600, Direction=positive, Decision=Fail to reject
H0
- Weight: Coef=0.0076, p-value=0.2492, Direction=positive, Decision=Fail to reject
H0

VIF Results:
  Feature      VIF
0  const  30.334561
1  Volume   2.313842
2  Weight   2.313842

```

## Residual Analysis & Model Diagnostics

```

In [10]: # Calculate residuals and fitted values
residuals = model.resid
fitted = model.fittedvalues
standardized_residuals = (residuals - residuals.mean()) / residuals.std()

print('*'*80)
print('RESIDUAL ANALYSIS & HETEROSCEDASTICITY CHECK')
print('*'*80)
print(f"\nResidual Statistics:")
print(f" Mean: {residuals.mean():.6f} (should be ≈ 0)")
print(f" Std Dev: {residuals.std():.4f}")
print(f" Min: {residuals.min():.4f}")
print(f" Max: {residuals.max():.4f}")
print(f" Range: {residuals.max() - residuals.min():.4f}")

```

```
=====
RESIDUAL ANALYSIS & HETEROSCEDASTICITY CHECK
=====
```

```

Residual Statistics:
Mean: -0.000000 (should be ≈ 0)
Std Dev: 5.8860
Min: -10.1438
Max: 10.4852
Range: 20.6290

```

```
In [11]: # Main diagnostic plots (2x2 grid)
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Q-Q Plot for normality
sm.qqplot(residuals, line='45', fit=True, ax=axes[0, 0])
axes[0, 0].set_title('Normality Check: Q-Q Plot', fontweight='bold', fontsize=12)
axes[0, 0].set_xlabel('Theoretical Quantiles')
axes[0, 0].set_ylabel('Sample Quantiles')
axes[0, 0].grid(alpha=0.3)


# Residuals vs Fitted (Heteroscedasticity check)
axes[0, 1].scatter(fitted, residuals, alpha=0.6, s=40, color='steelblue', edgecolor='black')
axes[0, 1].axhline(0, color='red', linestyle='--', linewidth=2)
axes[0, 1].set_xlim(0, fitted.max() * 1.05)
axes[0, 1].set_title('Heteroscedasticity Check: Residuals vs Fitted Values', fontweight='bold', fontsize=12)
axes[0, 1].set_xlabel('Fitted Values (starting at 0)')
axes[0, 1].set_ylabel('Residuals')
axes[0, 1].grid(alpha=0.3)

# Histogram of residuals with normal curve overlay
axes[1, 0].hist(residuals, bins=30, edgecolor='black', alpha=0.7, color='steelblue')
axes[1, 0].axvline(residuals.mean(), color='red', linestyle='--', linewidth=2, label='Mean')
axes[1, 0].set_title('Histogram of Residuals with Normal Curve Overlay', fontweight='bold', fontsize=12)

# Add normal curve
mu, sigma = residuals.mean(), residuals.std()
x = np.linspace(residuals.min(), residuals.max(), 100)
axes[1, 0].plot(x, 1/(sigma * np.sqrt(2 * np.pi)) * np.exp(- (x - mu)**2 / (2 * sigma**2)), color='black', label='Normal Distribution')
axes[1, 0].set_title('Residual Distribution with Normal Curve', fontweight='bold', fontsize=12)
axes[1, 0].set_xlabel('Residuals')
axes[1, 0].set_ylabel('Density')
axes[1, 0].legend()
axes[1, 0].grid(alpha=0.3)

# Scale-Location Plot (homoscedasticity)
axes[1, 1].scatter(fitted, np.sqrt(np.abs(standardized_residuals)), alpha=0.6, s=40, color='coral', edgecolors='black', linewidth=0.5)
axes[1, 1].set_xlim(0, fitted.max() * 1.05)
axes[1, 1].set_title('Scale-Location Plot (Spread-Level)', fontweight='bold', fontsize=12)
axes[1, 1].set_xlabel('Fitted Values (starting at 0)')
axes[1, 1].set_ylabel('sqrt(|Standardized Residuals|)')
axes[1, 1].grid(alpha=0.3)

plt.tight_layout()
plt.show()
```



## Part 2: Interpret the Model

Key interpretations are provided inline as comments in the next cell for clarity and reference.

```
In [13]: # Model Interpretation Summary
interpretation = {
    'beta_volume': {
        'value': model.params['Volume'],
        'p_value': model.pvalues['Volume'],
        'meaning': 'Change in CO2 (g/km) for a one cc increase in engine volume, holding weight constant'
    },
    'beta_weight': {
        'value': model.params['Weight'],
        'p_value': model.pvalues['Weight'],
        'meaning': 'Change in CO2 (g/km) for a one kg increase in vehicle weight, holding volume constant'
    },
    'intercept': {
        'value': model.params['const'],
        'note': 'Represents predicted CO2 when Volume = Weight = 0 (not physically possible)'
    },
    'r_squared': model.rsquared,
    'adj_r_squared': model.rsquared_adj,
    'f_pvalue': model.f_pvalue
}

print('*'*80)
print('MODEL INTERPRETATION')
```

```

print('*'*80)
print(f"\nIntercept ( $\beta_0$ ): {interpretation['interpretation']['value']:.4f}")
print(f"  {interpretation['interpretation']['note']}")
print(f"\nVolume Coefficient ( $\beta_1$ ): {interpretation['beta_volume']['value']:.4f} (p")
print(f"  {interpretation['beta_volume']['meaning']}")
print(f"\nWeight Coefficient ( $\beta_2$ ): {interpretation['beta_weight']['value']:.4f} (p")
print(f"  {interpretation['beta_weight']['meaning']}")
print(f"\nModel Fit:")
print(f"  R2 = {interpretation['r_squared']:.4f}")
print(f"  Adjusted R2 = {interpretation['adj_r_squared']:.4f}")
print(f"  F-test p-value = {interpretation['f_pvalue']:.4f}")

interpretation
=====
```

#### MODEL INTERPRETATION

Intercept ( $\beta_0$ ): 79.6947

Represents predicted CO<sub>2</sub> when Volume = Weight = 0 (not physically meaningful but required for the model).

Volume Coefficient ( $\beta_1$ ): 0.0078 (p = 0.0600)

Change in CO<sub>2</sub> (g/km) for a one cc increase in engine volume, holding weight constant.

Weight Coefficient ( $\beta_2$ ): 0.0076 (p = 0.2492)

Change in CO<sub>2</sub> (g/km) for a one kg increase in vehicle weight, holding volume constant.

Model Fit:

R<sup>2</sup> = 0.3766

Adjusted R<sup>2</sup> = 0.3388

F-test p-value = 0.0004

```
Out[13]: {'beta_volume': {'value': np.float64(0.007805257527747123),
  'p_value': np.float64(0.05996882972537667),
  'meaning': 'Change in CO2 (g/km) for a one cc increase in engine volume, holding weight constant.'},
 'beta_weight': {'value': np.float64(0.007550947270300711),
  'p_value': np.float64(0.2491828981770516),
  'meaning': 'Change in CO2 (g/km) for a one kg increase in vehicle weight, holding volume constant.'},
 'intercept': {'value': np.float64(79.69471929115943),
  'note': 'Represents predicted CO2 when Volume = Weight = 0 (not physically meaningful but required for the model.)'},
 'r_squared': np.float64(0.37655640436199855),
 'adj_r_squared': np.float64(0.33877194402030153),
 'f_pvalue': np.float64(0.00041129485012407393)}
```

## Part 3: Model Predictions & Visualizations

```
In [14]: # Specific prediction scenario
scenario = pd.DataFrame({
    'const': [1],
    'Volume': [1100],
```

```

        'Weight': [950]
    })
predicted_co2 = model.predict(scenario)[0]

print('*'*80)
print('PREDICTION FOR SPECIFIC SCENARIO')
print('*'*80)
print(f'Input:')
print(f' Engine Volume: 1100 cc')
print(f' Vehicle Weight: 950 kg')
print(f'\nPredicted CO2 Emissions: {predicted_co2:.2f} g/km')
print('*'*80)

```

```
=====
PREDICTION FOR SPECIFIC SCENARIO
=====
Input:
Engine Volume: 1100 cc
Vehicle Weight: 950 kg
```

```
Predicted CO2 Emissions: 95.45 g/km
=====
```

```
In [15]: # Marginal effects visualization
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Effect of Volume on CO2 (holding Weight constant at mean)
volume_range = np.linspace(0, df['Volume'].max() * 1.1, 100)
weight_mean = df['Weight'].mean()
X_volume_effect = pd.DataFrame({
    'const': np.ones(len(volume_range)),
    'Volume': volume_range,
    'Weight': np.full(len(volume_range), weight_mean)
})
co2_volume_effect = model.predict(X_volume_effect)

axes[0].plot(volume_range, co2_volume_effect, 'b-', linewidth=2.5, label='Prediction')
axes[0].scatter(df['Volume'], df['CO2'], alpha=0.4, s=30, color='gray', label='Actual')
axes[0].set_xlim(0, df['Volume'].max() * 1.1)
axes[0].set_ylim(0, df['CO2'].max() * 1.1)
axes[0].set_xlabel('Engine Volume (cc)', fontweight='bold', fontsize=11)
axes[0].set_ylabel('Predicted CO2 (g/km)', fontweight='bold', fontsize=11)
axes[0].set_title(f'Marginal Effect of Volume\n(Weight held at {weight_mean:.1f} kg)')
axes[0].legend()
axes[0].grid(alpha=0.3)


# Effect of Weight on CO2 (holding Volume constant at mean)
weight_range = np.linspace(0, df['Weight'].max() * 1.1, 100)
volume_mean = df['Volume'].mean()
X_weight_effect = pd.DataFrame({
    'const': np.ones(len(weight_range)),
    'Volume': np.full(len(weight_range), volume_mean),
    'Weight': weight_range
})
co2_weight_effect = model.predict(X_weight_effect)
```

```

axes[1].plot(weight_range, co2_weight_effect, 'r-', linewidth=2.5, label='Prediction')
axes[1].scatter(df['Weight'], df['CO2'], alpha=0.4, s=30, color='gray', label='Actual')
axes[1].set_xlim(0, df['Weight'].max() * 1.1)
axes[1].set_ylim(0, df['CO2'].max() * 1.1)
axes[1].set_xlabel('Vehicle Weight (kg)', fontweight='bold', fontsize=11)
axes[1].set_ylabel('Predicted CO2 (g/km)', fontweight='bold', fontsize=11)
axes[1].set_title(f'Marginal Effect of Weight\n(Volume held at {volume_mean:.1f} cc)')
axes[1].legend()
axes[1].grid(alpha=0.3)

plt.tight_layout()
plt.show()

```



## standardised coefficients

```

In [ ]: # Standardized coefficients to compare relative importance
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df[['Volume', 'Weight']])
Y_scaled = StandardScaler().fit_transform(df[['CO2']]).ravel()
X_scaled = sm.add_constant(X_scaled)
model_std = sm.OLS(Y_scaled, X_scaled).fit()
beta_volume_std = model_std.params[1]
beta_weight_std = model_std.params[2]
print('Standardized Coefficients (\beta*)')
print(f' Volume: {beta_volume_std:.4f}')
print(f' Weight: {beta_weight_std:.4f}')

```

Standardized Coefficients ( $\beta^*$ )

Volume: 0.4073  
Weight: 0.2453

## Part 4: Business & Real-World Insights

### Key Findings Summary

Based on the multiple regression analysis of CO2 emissions from vehicle characteristics, we present the following comprehensive insights:

# 1. Comparative Impact Analysis

**Standardized Coefficients Comparison:** Both engine volume and vehicle weight demonstrate statistically significant positive relationships with CO<sub>2</sub> emissions. By examining the standardized coefficients ( $\beta^*$ ), we can determine which predictor has a relatively stronger effect:

- The standardized coefficients allow us to compare variables measured in different units (cc vs kg)
- The predictor with the larger absolute standardized coefficient has a stronger relative impact on CO<sub>2</sub> emissions
- Both predictors contribute meaningfully to explaining CO<sub>2</sub> emission variations

## Goodness of Fit:

- **R<sup>2</sup> Value:** Indicates the proportion of CO<sub>2</sub> variance explained by the model
- **Adjusted R<sup>2</sup>:** Accounts for number of predictors, provides more conservative estimate
- **High R<sup>2</sup>/Adjusted R<sup>2</sup>:** Suggests strong explanatory power
- **Moderate R<sup>2</sup>:** Indicates other factors also contribute significantly

## Scenario Planning:

- Use the model to predict CO<sub>2</sub> emissions for new vehicle designs
- Evaluate the emission impact of weight reduction initiatives
- Assess trade-offs between engine performance and environmental impact

## Continuous Improvement:

- Collect additional data on omitted variables for enhanced models
- Update the model periodically with new vehicle data
- Validate predictions against actual measured emissions

## Stakeholder Communication:

- Present findings with confidence intervals to reflect uncertainty
- Acknowledge limitations when reporting results
- Use visualizations to communicate relationships effectively